

Self-Taught Semi-Self Speculative Decoding

S4D: a learned routing policy that **Pareto-dominates** hierarchical speculative decoding

viasd-rl

June 20, 2026

The new state of the art in hierarchical speculative decoding.

A strict Pareto improvement: up to $3.5\times$ **faster**, while *preserving* accuracy (lossless quality at $4\times$ **fewer** model calls).

We introduce the first reinforcement-learned routing policy for hierarchical speculative decoding, and set a new state of the art.

VIA-SD adds a slim intermediate verifier, but routes between tiers using confidence thresholds tuned by hand.

We present **three key contributions**:

1. A **learned per-token gating policy** that recasts tier-routing as a sequential decision problem, so a tiny q -free MLP replaces the fixed thresholds.
2. A **reward** that balances output quality against the latency of expensive verifier calls.
3. A **per-token RL formulation** with principled credit assignment (REINFORCE, then GRPO, then a per-token advantage).

The payoff: we match *lossless* accuracy (**0.91**) at **$4\times$ fewer** full-model calls, **$2.2\times$ faster** than greedy. We **Pareto-dominate** the paper's VIA-SD at roughly **$2\times$ its speedup** and **$2.3\times$ its accuracy**, and reach up to **$3.5\times$** when pushed for speed.

The bottleneck: LLM decoding is slow and sequential


- Generating each token streams the entire model from memory, so LLM inference is memory-bandwidth bound and dominated by sequential forward passes.
- Speculative decoding fixes this: a cheap *drafter* guesses several tokens, and the big *verifier* checks them all in one parallel pass.
- Accepted guesses give many tokens per expensive forward. It is **lossless**, and runs about **1.4×** faster (wall-clock) on our models.
- The one lever is the **acceptance rate**, because every rejected token forces a full-model step.

VIA-SD improves on vanilla SD

Vanilla SD makes a **binary** choice per token: accept the drafter, or reject and pay for the full model q .

VIA-SD adds a **slim middle verifier** q' (the target q with about 45% of its layers skipped), giving **three** choices. Borderline “middle-zone” tokens are handled cheaply by q' , so the full model is called far less often.

Since q' is the target verifying its *own* draft through a reduced copy of itself, this is the **Semi-Self** in S4D.



But it routes with thresholds that are picked by hand

VIA-SD makes the three-way choice with **two fixed confidence thresholds** on q' :

$$\alpha_1 = 0.5$$

$$\alpha_2 = 0.3.$$

- They are **global and static**: one cutoff for every token and every context.
- They are **brittle**, and must be re-tuned for each new model and task.
- They are **sub-optimal**: even after a full sweep, accuracy peaks at 0.50 on our setup.


Choosing a tier for each token, from its context, is a sequential decision problem. **So we learn it** — the gate teaches itself. That is the **Self-Taught** in S4D.

Contribution 1/3: a learned per-token gating policy (MDP)

We recast tier-routing as a **Markov decision process**:

- **State** s_t : q -free features (drafter/ q' confidence, agreement, $\text{KL}(p\|q')$, position).
- **Action** $a_t \in \{\text{accept, regenerate, escalate}\}$.
- **Policy** $\pi_\theta(a_t \mid s_t)$: a tiny MLP.

The full model q is consulted only to *score* at train time. It is never a gate input, so the gate is free to run at inference.



Contribution 2/3: a reward balancing quality and cost

Two ingredients per token:

$$\text{quality: } \text{match}_t = \mathbf{1}[\hat{y}_t = \arg \max_v q(v \mid x_{<t})]$$

$$\text{cost: } \text{cost}_t = (t_{p_1} + t_{q'}/\gamma + \mathbf{1}[\text{ESCALATE}] t_q) / t_{q_1}$$

Reward (cost in units of one full-model step):

$$\text{trajectory: } R = \overline{\text{match}} - \lambda \overline{\text{cost}} \text{ (+ } r_c \text{ correct)} \qquad \text{per-token: } r_t = \text{match}_t - \lambda \text{cost}_t$$

λ is the quality-versus-speed knob. `match` is a *dense* proxy for the sparse true objective (final-answer correctness).

Contribution 3/3: per-token RL with credit assignment

Policy gradient (all variants, plus an entropy bonus $\eta \mathcal{H}[\pi_\theta]$):


$$\nabla_\theta J = \mathbb{E} \left[\sum_t A_t \nabla_\theta \log \pi_\theta(a_t | s_t) \right]$$

They differ only in the **advantage** A_t :

- **REINFORCE** (trajectory): $A = R - b$, with EMA baseline $b \leftarrow (1-\tau)b + \tau R$.
- **GRPO** (group): $A_i = \frac{R_i - \text{mean}_{\mathcal{G}} R}{\text{std}_{\mathcal{G}} R}$, with PPO-clip $\rho_t = \frac{\pi_\theta}{\pi_{\theta_{\text{old}}}}$ and KL to π_{ref} .
- **per-token (v2)**: $\hat{A}_t = \frac{r_t - \text{mean}_{\text{batch}} r}{\text{std}_{\text{batch}} r}$, local credit and the **lowest variance**.

Training converges, and per-token (v2) is cleanest

RL training curves (0.5B→14B, GSM8K)




Results: the learned gate dominates

Method	GSM8K acc	q-calls/tok	speedup	
greedy- q (reference)	0.887	1.000	1.00×	
plain SD (standard, lossless)	0.907	0.259	1.40×	[†]
via_fixed (paper, tuned)	0.40	0.31	1.76×	[†] wall-clock; via_*
via_imit (imitation only)	0.34	0.29	1.69×	
via_rl REINFORCE	0.880	0.243	2.29×	
via_rl GRPO	0.853	0.162	2.94×	
via_rl REINFORCE+DIMR	0.907	0.250	2.22×	
via_rl v2 per-token	0.713	0.104	3.53×	

speedups are the overhead-free bandwidth model (not directly comparable).


⇒ **S4D is state-of-the-art hierarchical speculative decoding.**

Latency: speedup by method



plain SD = measured wall-clock; via_* = overhead-free bandwidth model.

The accuracy and speed frontier



Ablation: which model should we scale?

Setup	GSM8K acc	q-calls/tok	note
0.5B → 14B (main)	0.88–0.91	0.10–0.25	baseline pair
3B → 14B (bigger <i>drafter</i>)	0.80	0.028	36× fewer full-model calls
0.5B → 32B (bigger <i>verifier</i>)	0.33	0.07	accuracy collapses

- A stronger **drafter** raises acceptance, so the gate escalates far less at similar accuracy.
- A bigger **verifier** with the same weak drafter widens the gap, and the accept-heavy gate emits tokens the target would reject.

The lever is the drafter, not the verifier.

Ablation: the slim verifier q' and the cost knob λ

- **Better q' via DIMR layer-skip search:** REINFORCE 0.88 $\rightarrow$ 0.91. A trustworthy regenerate tier lets the gate use q' instead of escalating, and the tokens come out right.
- **But q' quality is moot once the gate is accept-heavy** (the v2 policy routes around the regenerate tier entirely).
- **The λ knob is a real frontier:** $\lambda=0.1 \rightarrow 0.80$, $0.2 \rightarrow 0.78$, $0.3 \rightarrow 0.71$ accuracy.
- **Honest negative:** optimizing final-answer correctness *directly* (instead of the match proxy) degraded every warm-start we tried; the match-trained REINFORCE+DIMR (0.91) stays best.

Impact: a new accuracy and efficiency frontier

- **Lossless quality, 4× cheaper.** We match greedy/lossless accuracy (**0.91**) at 4× **fewer** full-model calls, a 2.2× speedup, with *zero accuracy lost*.
- **A tunable frontier.** One knob gives 2.7× at −0.1 accuracy, or 3.5× at −0.19, up to 10× **fewer** full-model calls.
- **Pareto-dominates VIA-SD.** About 2× its speedup *and* +0.3 **to** +0.5 accuracy. Where its fixed thresholds collapse to 0.40, our gate holds **0.71 to 0.91**.
- **Scales with the drafter.** A 3B drafter gives 36× **fewer** full-model calls at 0.80 accuracy, and it is a free-to-deploy tiny MLP.

S4D is the new state of the art in hierarchical speculative decoding.

Stop hand-picking thresholds. Let the model schedule its own compute.

× = bandwidth model; “fewer calls” (q-calls/token) is exact and hardware-independent.

Future work

- **Spec-spec decoding:** make the drafter itself speculative, a recursive hierarchy where the gate routes across *levels*.
- **KnapSpec:** treat verification as a **knapsack**, spending a fixed budget of expensive calls on the highest-value tokens (value is correctness impact, weight is cost).
- **Correctness-floor RL:** minimize cost *subject to* an accuracy floor, which avoids the reward-hacking collapse.
- **Scale:** bigger drafters, more tasks, and real-engine wall-clock (static cache, CUDA graphs).

S4D

Self-Taught Semi-Self Speculative Decoding

The new state of the art in hierarchical speculative decoding.

Stop hand-picking thresholds. Let the model schedule its own compute.

Thanks! questions?